

# Yash Bheke

linkedin.com/in/yash-bheke|github.com/yashbheke2000|yash.bheke2000@gmail.com|yash-bheke.lovable.app

## SUMMARY

Data scientist and ML-focused graduate student with end-to-end experience building scalable data pipelines, accelerating model inference by 87%, and driving six-figure cost savings through production ML systems, A/B testing, and cloud-based analytics on Databricks, AWS, and MLOps stacks. Passionate about building production-ready solutions that drive measurable business impact.

## EDUCATION

**Master of Science in Information Systems** - University of Cincinnati Aug 2024 - Apr 2026

Courses: Gen AI, Statistical Computing, Data Mining for BI, Data Analysis, AI & Machine Learning, Data Visualization

**Bachelor of Engineering in Electronics and Telecommunication** - Mumbai University Aug 2018 - May 2022

## TECHNICAL SKILLS

**Languages:** Python, SQL, R

**Machine Learning & Statistics:** TensorFlow, Scikit-learn, PyTorch, NLP, A/B Testing

**Big Data & Cloud:** Databricks, PySpark, Snowflake, AWS, BigQuery

**Visualization & Analysis:** Tableau, Power BI, Looker, Matplotlib

**Deployment:** Docker, FastAPI, Git, GitHub, CI/CD

## WORK EXPERIENCE

**Parking Base** (Data Science Intern) Feb 2026 - Apr 2026

**Tech Stack:** Looker, BigQuery, Gemini SDK, Python, LookML

- Developed a **vendor-level profitability matrix** leveraging BigQuery to analyze **450K+ transactions** to quantify key margin drivers across booking channels, identifying commission misalignments contributing to **8.3% leakage in profit margins**.
- Reduced reconciliation review time by conducting anomaly detection analysis on payroll discrepancies (biometric vs. manual clock-outs), detecting outlier cases with **>15% time variance** and improving payroll review controls.
- Raised payment gateway approval rate to **98.7%** (+5.7% relative) by building a statistical segmentation of payment failure patterns that surfaced declines by card type, reason and channel, flagging high failure rate methods for review.

**University of Cincinnati** (AI Engineer Intern) Sept 2025 - Dec 2025

**Tech Stack:** Python, Tensorflow, Keras, N8N, RAG, vector store, LLM

- Developed a **30-intent text classification model** for University admissions chatbot using TensorFlow/Keras on 1.4K+ labeled queries; improved weighted F1-score from 0.82 to 0.91 through hyperparameter tuning and class rebalancing.
- Automated ingestion and embedding updates for new FAQ documents, eliminating the need for full model retraining.

**University of Cincinnati** (Graduate Assistant) Sept 2025 - Dec 2025

**Tech Stack:** Databricks, DBeaver, ETL, Python, PySpark, SQL, AWS, S3, RDS

- Optimized large-scale data transformation workflows in Databricks (PySpark/SQL) processing **15M+ records**, reducing end-to-end runtime by **31%** and improving data availability for faster downstream analysis.
- Built analytics-ready datasets by integrating AWS RDS and S3 sources; standardized schemas, implemented data validation checks and engineered derived features to support statistical modeling and ML experiments.

**Accelya** (Data Scientist) Sept 2022 - Jul 2024

**Tech Stack:** Oracle, Python, SQL, ETL, A/B Testing, Fast API, Process Automation

- Accelerated model inference time by **87% (from 15+ minutes to under 2 minutes)** for the predictive analytics pipeline with PCA-based dimensionality reduction and XGBoost hyperparameter tuning which aided business analysts across 20+ clients.
- Saved **\$370K+ annually** in operational costs by developing and deploying ensemble classification models (Random Forest, Gradient Boosting) via FastAPI, achieving **83% accuracy** in flight delay predictions for 12 clients with **1k+ daily inferences**.
- Improved user engagement by **18%** and reduced reporting errors by **16%** by designing A/B testing frameworks that combined Bayesian inference and hypothesis testing; identified **9 statistically significant improvements** in financial workflows and UI/UX.
- Automated data pre-processing, eliminating **60+ hours** of monthly manual effort by streamlining cleaning, validation, and EDA.

## PROJECTS

**Airbnb Optimum Pricing Tool** [ DataBricks, Python, sklearn, SHAP ] Oct 2024 - Dec 2024

- Developed a regression model to predict listing prices, improving R<sup>2</sup> by **35.8%** (0.41 to 0.56) compared to baseline models.
- Performed data cleaning, pre-processing, and feature engineering (outlier removal, imputation, one-hot encoding), applied SHAP to identify key price drivers and deployed via Streamlit for hosts to tweak parameters and view optimum pricing.

## EXTRACURRICULAR ACTIVITIES

**Project Lead, Neo Initiative (Pro-bono Consulting)**, University of Cincinnati Sept 2025 - Dec 2025

**Student Ambassador, MSIS**, University of Cincinnati Sept 2024 - Aug 2025